

✓ CSCA-5028: Application of Software Architecture for Big Data

- Author: Alexander Meau
- Email: alme9155@colorado.edu
- GitHub: <https://github.com/alme9155/csca-5028-sentiment-analysis>

✓ I. Brief description of the problem and data

Project Name: AI-Powered Movie Sentiment Rating System

This project aims to fine-tune a pre-trained NLP model and adapt the natural language processor to provide proper ratings based on review text. This notebook offers a development playground for the software project.

- **What problem is the product aimed at solving?**

- This product aims to address the challenges of generating accurate rating scores for movie review text.

- **Who is the product geared towards (targeted audience)?**

- The primary target audience is online movie platforms, such as Netflix or Amazon Prime, to use these ratings in their recommender system.

- **How is the product unique?**

- This product is unique in its application of transfer learning, where a pre-trained NLP model is fine-tuned for the domain of movie reviews.

Dataset:

- Stanford Sentiment Treebank Class 5(SST-5):

- SST-5 is a fine-grained sentiment analysis dataset derived from Rotten Tomatoes movie review snippets

created by Stanford NLP Laboratory.

- This is a fine-grained sentiment analysis dataset because the sentiment labels are broken down to 5 rating categories. Most of the sentiment dataset only has classification between 0 and 1 (positive/negative).
- Rating categories range from 0-5: very negative, negative, neutral, positive, or very positive.

Data Size and Dimension

- The dataset is organized into 3 splits: *Train*, *Validation*, *Test*.
- The *Train* split contains 8544 sentence samples.
- Each sample includes:
 - "text": The review text
 - "label": Numeric sentiment class (0-5)
 - "label_text": Corresponding sentiment class label (e.g. very positive).
- Example: "This is a great movie", 4, "very positive"

```
# Step 1: Install PyTorch with CUDA 11.8 (required for GPU), and latest Hugging Face system
!pip install -q torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118
!pip install -q -U "transformers>=4.44.0" "datasets>=3.0.0" "accelerate>=0.33.0" "evaluate>=0.4.0"

import torch
print(f"PyTorch: {torch.__version__}")
print(f"CUDA available: {torch.cuda.is_available()}")
print(f"GPU: {torch.cuda.get_device_name(0)} if torch.cuda.is_available() else 'None'")
```

```
511.6/511.6 kB 18.9 MB/s eta 0:00:00
84.1/84.1 kB 8.5 MB/s eta 0:00:00
9.5/9.5 MB 107.5 MB/s eta 0:00:00
47.7/47.7 MB 43.2 MB/s eta 0:00:00
```

PyTorch: 2.8.0+cu126
CUDA available: True
GPU: Tesla T4

✗ II. Exploratory Data Analysis (EDA)

- Examine video quality:
 - Dimension (width x height)
 - Frame Count and Frame per secound
 - Total count of video clips
 - Total count of video category

```
from google.colab import userdata
userdata.get('HF_TOKEN')

from datasets import load_dataset
ds = load_dataset("SetFit/sst5")
ds
```

```
Repo card metadata block was not found. Setting CardData to empty.
WARNING:huggingface_hub.repocard:Repo card metadata block was not found. Setting CardData to empty.
DatasetDict({
    train: Dataset({
        features: ['text', 'label', 'label_text'],
        num_rows: 8544
    })
    validation: Dataset({
        features: ['text', 'label', 'label_text'],
        num_rows: 1101
    })
    test: Dataset({
        features: ['text', 'label', 'label_text'],
        num_rows: 2210
    })
})
```

```
from datasets import load_dataset
import pandas as pd
from IPython.display import display
```

```
print("===== DATASET DIMENSIONS =====\n")
print(f"Number of rows (samples): {len(ds["train"])}")
print(f"Number of columns (features): {len(ds["train"].features)}\n")

print("Feature schema:")
print(ds["train"].features)

#df_train_head = ds["train"].to_pandas().head(10)
#display(df_train_head)

# Show a few training samples
for i in range(3):
    print(f"\nExample {i}:")
    print("Text :", ds["train"][i]["text"])
    print("Label:", ds["train"][i]["label"])

===== DATASET DIMENSIONS =====

Number of rows (samples): 8544
Number of columns (features): 3

Feature schema:
{'text': Value('string'), 'label': Value('int64'), 'label_text': Value('string')}

Example 0:
Text : a stirring , funny and finally transporting re-imagining of beauty and the beast and 1930s h
Label: 4

Example 1:
Text : apparently reassembled from the cutting-room floor of any given daytime soap .
Label: 1


Example 2:
Text : they presume their audience wo n't sit still for a sociology lesson , however entertainingly
Label: 1
```


```
import matplotlib.pyplot as plt
import numpy as np
```

```
train_txt_length = [len(ex["text"]) for ex in ds["train"]]

plt.hist(train_txt_length, bins=30)
plt.xlabel("Text length (characters)")
plt.ylabel("Frequency")
plt.title("SST-5 Train Split: Text Length Distribution")
plt.show()

train_arr = np.array(train_txt_length)
print(f"\n\"Train\" split \"text\" length (characters):")
print(f"  mean: {train_arr.mean():.1f}")
print(f"  std : {train_arr.std():.1f}")
print(f"  min : {train_arr.min()}")
print(f"  max : {train_arr.max()}")
for q in [25, 50, 75, 90, 95]:
    print(f"  {q}th percentile: {np.percentile(train_arr, q):.1f})")
```






```
"Train" split "text" length (characters):
mean: 102.7
std : 51.7
min : 4
max : 283
25th percentile: 62.0
50th percentile: 98.0
75th percentile: 137.0
90th percentile: 173.0
95th percentile: 194.0
```

```
import matplotlib.pyplot as plt
from collections import Counter

train_labels = [ex["label_text"] for ex in ds["train"]]
label_counts = Counter(train_labels)
ordered_labels = ["very negative", "negative", "neutral", "positive", "very positive"]
counts = [label_counts[label] for label in ordered_labels]

# --- Bar chart ---
plt.figure(figsize=(6,4))
plt.bar(ordered_labels, counts)
plt.xlabel("Sentiment label")
plt.ylabel("Frequency")
plt.title("SST-5 Train Split: label_text Distribution")
plt.xticks(rotation=30)
plt.tight_layout()
plt.show()

print("\n" "Train" split label_text distribution:")
for label, count in zip(ordered_labels, counts):
    print(f" {label:13s}: {count}")
```



"Train" split label_text distribution:

```
very negative: 1092
negative      : 2218
neutral       : 1624
positive      : 2322
very positive: 1288
```

III. Fine-tuning DistilBERT NLP model for movie review rating classification

For this project, we will download a pre-trained natural language processing model and then fine-tune it to adapt to

movie review ratings using the Stanford SST-5 dataset..

- **Model Name:** distilbert/distilbert-base-uncased
- **Params:** Trained with 67M parameters
- **Size:** ~268 MB
- **Use Cases:** Text classification, sentiment analysis.
- **Tokenizer:** WordPiece, vocab size ~30K lower-case only(uncased)
- **HuggingFace:** <https://huggingface.co/distilbert/distilbert-base-uncased>

```
from transformers import AutoTokenizer, DataCollatorWithPadding
import numpy as np, evaluate

model_name = "distilbert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)

# Even though max length of 'text' field is 255
# number of converted tokens should be much less than 128
def preprocess(examples):
    return tokenizer(
        examples["text"],
        truncation=True,
        max_length=128,
    )

tokenized_ds = ds.map(preprocess, batched=True)
tokenized_ds = tokenized_ds.rename_column("label", "labels")
tokenized_ds.set_format(type="torch")
#, columns=["input_ids", "attention_mask", "labels"])

data_collator = DataCollatorWithPadding(tokenizer)

accuracy_metric = evaluate.load("accuracy")
f1_metric = evaluate.load("f1")

def compute_metrics(eval_pred):
    logits, labels = eval_pred
```

```
logits, labels = eval_pred
preds = np.argmax(logits, axis=-1)
return {
    "accuracy": accuracy_metric.compute(predictions=preds, references=labels)["accuracy"],
    "f1": f1_metric.compute(predictions=preds, references=labels, average="weighted")["f1"],
}
```

tokenizer_config.json: 100%	48.0/48.0 [00:00<00:00, 6.29kB/s]
config.json: 100%	483/483 [00:00<00:00, 62.3kB/s]
vocab.txt: 100%	232k/232k [00:00<00:00, 1.44MB/s]
tokenizer.json: 100%	466k/466k [00:00<00:00, 2.85MB/s]
Map: 100%	8544/8544 [00:00<00:00, 28634.41 examples/s]
Map: 100%	1101/1101 [00:00<00:00, 20344.55 examples/s]
Map: 100%	2210/2210 [00:00<00:00, 24365.75 examples/s]
Downloading builder script:	4.20k/? [00:00<00:00, 483kB/s]
Downloading builder script:	6.79k/? [00:00<00:00, 778kB/s]

```
from transformers import AutoModelForSequenceClassification

pairs = {}
for lbl, lbl_txt in zip(ds["train"]["label"], ds["train"]["label_text"]):
    if lbl not in pairs:
        pairs[int(lbl)] = lbl_txt

# Sort by numeric label ID just to be safe
id2label = {k: pairs[k] for k in sorted(pairs.keys())}
label2id = {v: k for k, v in id2label.items()}

model = AutoModelForSequenceClassification.from_pretrained(
    model_name
```

```
model.c_name,
num_labels=5,
id2label=id2label,
label2id=label2id,
)
```

model.safetensors: 100%

268M/268M [00:01<00:00, 151MB/s]

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and

```
import transformers
from transformers import TrainingArguments, Trainer
print(transformers.__version__)

training_args = TrainingArguments(
    output_dir=".sst2-distilbert",
    eval_strategy="epoch",
    save_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=32,
    per_device_eval_batch_size=32,
    num_train_epochs=3,
    weight_decay=0.01,
    load_best_model_at_end=True,
    metric_for_best_model="accuracy",
    push_to_hub=False,
    logging_dir='./logs',
    report_to=[],
)
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_ds["train"],
    eval_dataset=tokenized_ds["validation"],
```

```
        tokenizer=tokenizer,  
        data_collator=data_collator,  
        compute_metrics=compute_metrics,  
    )  
    trainer.train()
```

4.57.1

```
/tmp/ipython-input-2210952062.py:21: FutureWarning: `tokenizer` is deprecated and will be removed in  
    trainer = Trainer(  
[801/801 01:53, Epoch 3/3]
```

Epoch	Training Loss	Validation Loss	Accuracy	F1
1	No log	1.174383	0.485014	0.452012
2	1.181200	1.159466	0.482289	0.467989
3	1.181200	1.169203	0.484105	0.473761


1	No log	1.174383	0.485014	0.452012
2	1.181200	1.159466	0.482289	0.467989
3	1.181200	1.169203	0.484105	0.473761

```
TrainOutput(global_step=801, training_loss=1.0918719307164872, metrics={'train_runtime': 115.0378,  
'train_samples_per_second': 222.814, 'train_steps_per_second': 6.963, 'total_flos':  
339896411838720.0, 'train_loss': 1.0918719307164872, 'epoch': 3.0})
```

```
# Plot training loss vs epoch  
import matplotlib.pyplot as plt  
  
history = trainer.state.log_history  
log_df = pd.DataFrame(history)  
display(log_df.head())  
  
eval_df = log_df[log_df["eval_accuracy"].notnull()]  
  
plt.figure()  
plt.plot(eval_df["epoch"], eval_df["eval_accuracy"], marker="o", label="Accuracy")  
plt.plot(eval_df["epoch"], eval_df["eval_f1"], marker="s", label="F1 (weighted)")  
plt.xlabel("Epoch")  
plt.ylabel("Score")  
plt.title("Validation Accuracy & F1 per Epoch (SST-5)")  
plt.legend()
```

```
plt.legend()  
plt.grid(True)  
plt.show()
```

	eval_loss	eval_accuracy	eval_f1	eval_runtime	eval_samples_per_second	eval_steps_per_second
0	1.174383	0.485014	0.452012	1.4861	740.843	23.551
1	NaN	NaN	NaN	NaN	NaN	NaN
2	1.159466	0.482289	0.467989	1.4749	746.467	23.730
3	1.169203	0.484105	0.473761	1.5614	705.141	22.416
4	NaN	NaN	NaN	NaN	NaN	NaN



IV. Validate with accuracy measurement

- Use the fine-tuned DistilBERT model to evaluate on the **validation** and **test** splits of SST-5.
- Call `trainer.evaluate()` to compute key metrics, including **accuracy**, **F1 (weighted)**, and **loss**.
- Compare **validation vs. test accuracy** to check how well the model generalizes beyond the training data.
- This comparison result can be used to determine if re-training is required.

```
validation_results = trainer.evaluate(tokenized_ds["test"])
print("Validation set results:")
print("-----\n")
for k, v in validation_results.items():
    print(f"{k}: {v:.4f}" if isinstance(v, float) else f"{k}: {v}")
print("\n")

test_results = trainer.evaluate(tokenized_ds["test"])
print("Test set results:")
print("-----\n")
for k, v in test_results.items():
    print(f"{k}: {v:.4f}" if isinstance(v, float) else f"{k}: {v}")
print("\n")
```

[70/70 03:09]

Validation set results:

```
-----  
  
eval_loss: 1.1309  
eval_accuracy: 0.5100  
eval_f1: 0.4762  
eval_runtime: 3.1607  
eval_samples_per_second: 699.2140  
eval_steps_per_second: 22.1470  
----- 2.0000
```

```
epoch: 3.0000
```


```
Test set results:
```

```
-----  
eval_loss: 1.1309  
eval_accuracy: 0.5100  
eval_f1: 0.4762  
eval_runtime: 3.1048  
eval_samples_per_second: 711.7940  
eval_steps_per_second: 22.5460  
epoch: 3.0000
```

```
import matplotlib.pyplot as plt  
import numpy as np  
  
# --- Build comparison chart (validation vs test) ---  
metric_keys = ["eval_accuracy", "eval_f1", "eval_loss"]  
val_values = [validation_results[m] for m in metric_keys]  
test_values = [test_results[m] for m in metric_keys]  
  
x = np.arange(len(metric_keys))  
width = 0.35  
  
plt.figure()  
plt.bar(x - width/2, val_values, width, label="Validation")  
plt.bar(x + width/2, test_values, width, label="Test")  
  
plt.xticks(x, metric_keys, rotation=20)  
plt.ylabel("Score")  
plt.title("Validation vs Test Performance (SST-5)")  
plt.legend()  
plt.tight_layout()  
  
for i, v in enumerate(val_values):
```

```
plt.text(x[i] - width/2, v, f"{v:.3f}", ha="center", va="bottom")
for i, v in enumerate(test_values):
    plt.text(x[i] + width/2, v, f"{v:.3f}", ha="center", va="bottom")

plt.show()
```



▼ V. Save model and tokenizer

- Save the **fine-tuned model weights** and configuration to be used for production inference.

- Save the tokenizer files to the same directory so inference code can encode text exactly as during training.
- Keep track of the model version and training parameters for reproducibility and future updates.

```
# save model for online dockerization

save_dir = "./distilbert-sst5-finetuned"
trainer.save_model(save_dir)
tokenizer.save_pretrained(save_dir)
print(f"Saved model and tokenizer to: {save_dir}")
```

```
Saved model and tokenizer to: ./distilbert-sst5-finetuned
```

Vi. Reference:

- Model page: <https://huggingface.co/samadpls/sentiment-analysis>
- DataSet: <https://huggingface.co/datasets/SetFit/sst5>
- Source: <https://nlp.stanford.edu/sentiment/>